

What is Mutual Information and How it is used in Feature Selection?

Mutual information is a non-parametric metric that quantifies the dependency between two variables. It measures how much information about one variable can be obtained from another. In the context of feature selection for regression, mutual information can help identify both linear and non-linear relationships between features and the target variable. This scoring function is useful when the relationship is not strictly linear and may involve more complex dependencies.

Here is the code snippet:

```
from sklearn.feature_selection import SelectKBest, mutual_info_regression

# Create a SelectKBest object using mutual information
k_best = SelectKBest(score_func=mutual_info_regression, k=5) # Select the top 5 features
```

Mutual information (MI) is a mathematical measure used to quantify the amount of information shared between two random variables. In the context of feature selection, it is commonly employed to measure the dependence or information content shared between a feature (predictor variable) and a target variable (the variable to be predicted). MI is particularly useful when dealing with both linear and non-linear relationships between variables.

The mathematical calculation of mutual information is based on probability theory and information theory. For two discrete random variables X and Y, the mutual information (MI(X, Y)) is calculated as follows:

$$MI(X, Y) = \sum \sum p(x, y) * \log [p(x, y) / (p(x) * p(y))]$$

Where

- $p(x, y)$ is the joint probability distribution of X and Y (the probability that $X=x$ and $Y=y$).
- $p(x)$ is the marginal probability distribution of X (the probability that $X=x$, regardless of Y).
- $p(y)$ is the marginal probability distribution of Y (the probability that $Y=y$, regardless of X).
- The summation Σ is taken over all possible values of X and Y .

Here's a simplified step-by-step explanation of the calculation:

1. Calculate the joint probability distribution $p(x, y)$ for all possible combinations of values for X and Y in your dataset. This represents how often X and Y occur together.
2. Calculate the marginal probability distributions $p(x)$ and $p(y)$ for X and Y individually. These represent the probabilities of X and Y occurring independently of each other.
3. For each combination of X and Y , calculate the ratio $p(x, y) / (p(x) * p(y))$ and take the logarithm of this ratio.
4. Multiply the result from step 3 by $p(x, y)$ for each combination of X and Y and sum all these values over all combinations.

The resulting value, $MI(X, Y)$, provides a measure of how much information about X can be obtained from observing Y and vice versa. When used in feature selection, mutual information helps identify features that contain information that is relevant to the target variable, regardless of whether the relationship is linear or non-linear.

It's important to note that the calculation of mutual information can also be adapted for continuous random variables by using probability density functions and integrals instead of discrete probabilities and summations. The formula remains conceptually similar but may involve integration over the continuous variable space.